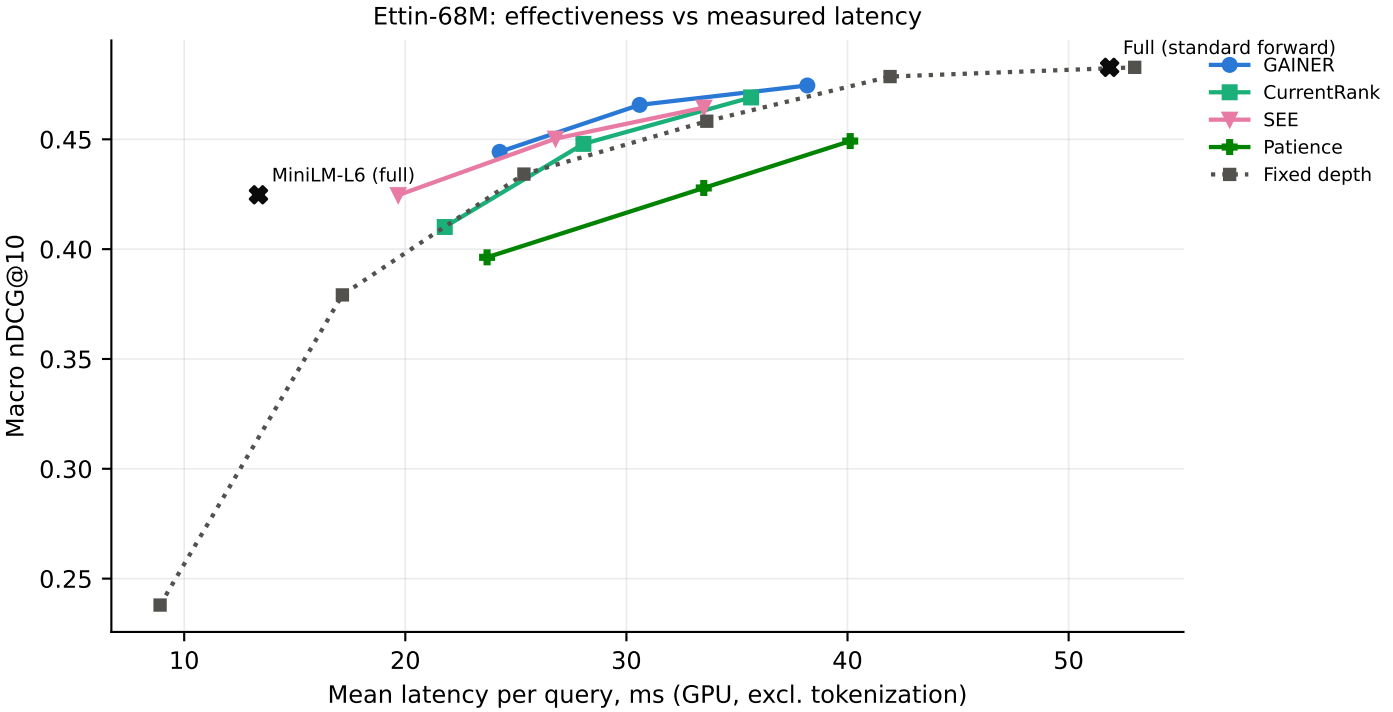
Ettin-68M: effectiveness vs measured latency
Macro nDCG@10
0.45
0.40
0.35
0.30
0.25
10
20
30
40
50
Mean latency per query, ms (GPU, excl. tokenization)
MiniLM-L6 (full)
Full (standard forward)
GAINER
CurrentRank
SEE
Patience
Fixed depth