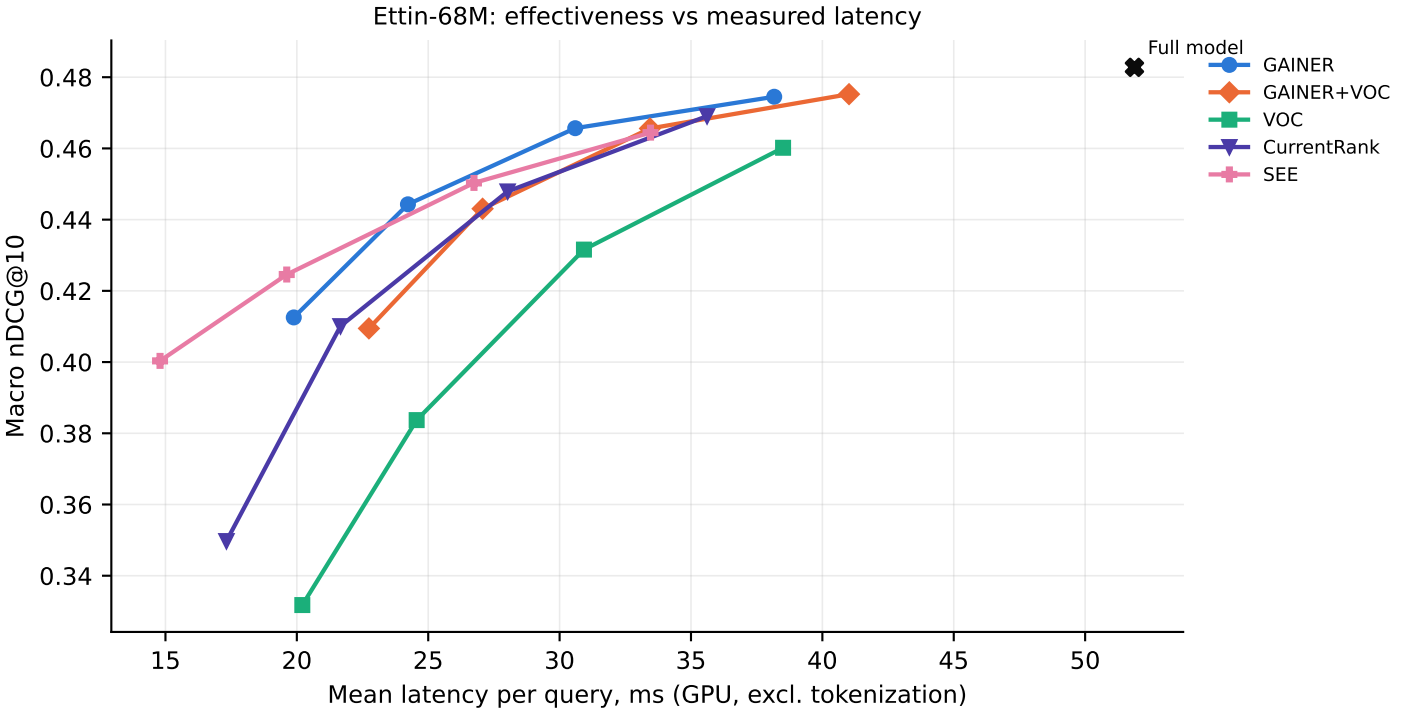
Ettin-68M: effectiveness vs measured latency
Full model
GAINER
GAINER+VOC
VOC
CurrentRank
SEE
Macro nDCG@10
0.48
0.46
0.44
0.42
0.40
0.38
0.36
0.34
15
20
25
30
35
40
45
50
Mean latency per query, ms (GPU, excl. tokenization)